


BOOM v2


an open-source out-of-order RISC-V core

Christopher Celio, Pi-Feng Chiu, Borivoje Nikolic, David Patterson, Krste Asanović
2017 RISC-V Workshop #7




Berkeley
Architecture
Research

What is BOOM?




- "Berkeley Out-of-Order Machine"
- out-of-order
- superscalar
- implements **RV64G**, boots Linux
- It is synthesizable
- it is open-source
- written in **Chisel** (16k loc)
- It is parameterizable generator
- built on top of Rocket-chip SoC Ecosystem

BOOM fits into Rocket-chip SoC




Lots of neat features

- advanced branch prediction
 - BTB, RAS (1-cycle)
 - gshare and TAGE-based conditional predictor implementations (3-cycle)
 - fits in single-port SRAM
- loads can issue out-of-order wrt to other loads, stores
- hardware performance counters (~40 events)



Parameterized Superscalar

dual-issue (5r,3w)




```

val exe_units = ArrayBuffer[ExecutionUnit]()
exe_units += Module(new ALUExeUnit(is_branch_unit      = true,
                                    , has_fpu          = true,
                                    , has_mul          = true
                                   ))
exe_units += Module(new ALUMemExeUnit(fp_mem_support = true,
                                       , has_div          = true
                                      ))

```


Parameterized Superscalar

dual-issue (5r,3w)



```
val exe_units = ArrayBuffer[ExecutionUnit]()
exe_units += Module(new ALUExeUnit(is_branch_unit = true,
                                    has_fpu = true,
                                    has_mul = true))
exe_units += Module(new ALUMemExeUnit(fp_mem_support = true,
                                      has_div = true))
```

Quad-issue (9r,4w)




OR

```
exe_units += Module(new ALUExeUnit(is_branch_unit = true))
exe_units += Module(new ALUExeUnit(has_fpu = true,
                                    has_mul = true))
exe_units += Module(new ALUExeUnit(has_div = true))
exe_units += Module(new MemExeUnit())
```

Good news




- BOOM has (finally) been taped out!
 - 2017 Aug 15
- TSMC 28 nm HPM
- Academic test-chip
- BOOM is in support of studying other features



BOOM v2

taping out an out-of-order processor
with a 2 person team in 4 months

Christopher Celio, Pi-Feng Chiu, Borivoje Nikolic, David Patterson, Krste Asanović
2017 CARRV Workshop



Berkeley
Architecture
Research



Comparison: Industry

Processor	SiFive U54 Rocket (RV64GC)	Berkeley BOOMv2	UltraSPARC T2	ARM Cortex-A9	Intel Xeon Ivy
Language	Chisel	Chisel	Verilog	-	SystemVerilog
Core LoC	8,000	16,000	290,000	-	NDA
Total LoC	34,000	50,000	1,300,000	-	NDA

Comparison: Industry

Processor	SiFive U54 Rocket (RV64GC)	Berkeley BOOMv2	UltraSPARC T2	ARM Cortex-A9	Intel Xeon Ivy
Language	Chisel	Chisel	Verilog	-	SystemVerilog
Core LoC	8,000	16,000	290,000	-	NDA
Total LoC	34,000	50,000	1,300,000	-	NDA
Foundry	TSMC	TSMC	TI	TSMC	Intel
Technology	28 nm (HPC)	28 nm (HPM)	65 nm	40 nm (G)	22 nm
Core+L1 Area	0.54 mm ²	0.52 mm ² 16kB/16kB	~12 mm ²	~2.5 mm ²	~12 mm ² core+L1+L2
Coremark/MHz	2.75	3.92	1.64*	3.71	5.60
Frequency	1.5 GHz	1.2 GHz**	1.17 GHz	1.4 GHz	3.3 GHz

*From eembc.org. 32 threads/8 cores achieve 13 Cm/MHz.

**Estimated



Limits of Educational Libraries


■ BOOMv1

- Educational 45 nm libraries
- CACTI model for SRAMs
- synthesized flip-flops for most non-cache memory arrays
- good for catching small timing bugs
- not accurate enough to justify sweeping redesigns

■ BOOMv2


- TSMC 28 nm HPM (high performance mobile)
- LVT-based standard cells
- memory compiler (one- and two-port)

4 months of agile tape-out




*ignore the Y-axis:


- too many parameters/variables changing between each run
- doesn't capture DRC violations




- short pipeline
 - inspired by R10K, 21264, Cortex-A9
- unified issue window




BOOMv1: Frontend



BOOMv1: Frontend




Frontend



BOOMv1



- BTB in SRAM
 - set-associative
 - partially tagged
 - Checker to verify integrity
- BPD (Conditional Predictor)
 - hash gets entire stage
 - redirect based on BTB
 - redirect pushed back (I\$)




BOOMv2


Regfile: (the first P&R)

- Regfile blows up
- critical paths in issue-select, register read






- split unified issue window (1 day)
- split physical register file (1 week)
- issue/register read separate stages (~1 day)
- 2 stage rename (~1 week)
- 3 stage fetch (3 weeks)




Regfile Challenges

- foundational IP provide memories with single and dual-port memories
- companies build their own hand-crafted register files
- not enough labor in academia to support a customized register file






synthesized




array



- Routing 4480 wires causes congestion issues

- Reduce read_data routing wires to 384

Regfile: Our quick solution




- hand-craft our own bit block out of standard cells
 - tri-state to drive read wires
 - hierarchical bitlines
- pre-place arrays of bit blocks
- let router handle the 18 wires per 1 flip-flop
- blackbox in Chisel
 - simulate control using Chisel model
 - verify at gate-level


Before and After



	BOOMv1	BOOMv2
BTB entries	40 (fully-associative)	64 x 4 (set-associative)
Fetch Width	2 insts	2 insts
Issue Width	3 micro-ops	4 micro-ops
Issue Entries	20	16/20/10
Regfile	7r3w	6r3w (int), 3r2w (fp)
	iALU+iMul+FMA	iALU+iMul+iDiv
Exe Units	iALU+fDiv Load/Store	iALU FMA+fDiv Load/Store



BOOMv1-2f3i



BOOMv2-2f4i

Results





Results

- Hard to compute -- not apples to oranges
 - a lot of the work was about design rule checks, not performance



Results

- Hard to compute -- not apples to oranges
 - a lot of the work was about design rule checks, not performance
- Roughly 25% decrease in clock period
- Roughly 20% decrease in Coremark/MHz
 - known issue from load-use delay

Results



- Hard to compute -- not apples to oranges
 - a lot of the work was about design rule checks, not performance
- Roughly 25% decrease in clock period
- Roughly 20% decrease in Coremark/MHz
 - known issue from load-use delay
- Barely a win?
- Changes fixed DRC errors, geometry errors (shorts),
SO...

Results



- Hard to compute -- not apples to oranges
 - a lot of the work was about design rule checks, not performance
- Roughly 25% decrease in clock period
- Roughly 20% decrease in Coremark/MHz
 - known issue from load-use delay
- Barely a win?
- Changes fixed DRC errors, geometry errors (shorts),
SO...
 - infinitely Faster! Better! Stronger!



Agile Hardware Development

- RTL hacking can be very agile
 - ~6 minutes to compile, build, and run "riscv-tests" regression suite (10 KHz for Verilator simulator)
 - Chisel allows for quick, far-reaching changes
 - generator approach allows for late-binding design decisions
 - small changes, improvements (that don't affect floor plan) are agile

Agile Hardware Development



- RTL hacking can be very agile
 - ~6 minutes to compile, build, and run "riscv-tests" regression suite (10 KHz for Verilator simulator)
 - Chisel allows for quick, far-reaching changes
 - generator approach allows for late-binding design decisions
 - small changes, improvements (that don't affect floor plan) are agile
- Physical design is a bottleneck
 - 2-3 hours for synthesis results
 - 8-24 hours for p&r results
 - too much value provided by manual intervention
 - reports are difficult to reason about
 - sadly can't throw this at a computer and get 7 good designs a week later

Agile Hardware Development

- RTL hacking can be very agile
 - ~6 minutes to compile, build, and run "riscv-tests" regression suite (10 KHz for Verilator simulator)
 - Chisel allows for quick, far-reaching changes
 - generator approach allows for late-binding design decisions
 - small changes, improvements (that don't affect floor plan) are agile
- Physical design is a bottleneck
 - 2-3 hours for synthesis results
 - 8-24 hours for p&r results
 - too much value provided by manual intervention
 - reports are difficult to reason about
 - sadly can't throw this at a computer and get 7 good designs a week later
- Verification is another bottleneck
 - I can write bugs faster than I can find them



Future Directions (for BOOM)

- A lot of improvements to make to BOOM
 - Clear direction of further IPC/QoR improvements

Future Directions (for me)



Esperanto Technologies

"Esperanto Technologies designs high-performance energy-efficient computing solutions."
- riscv.org

- See Dave Ditzel's talk from yesterday to learn more
- RISC-V Founding Gold Member
 - on lots of working groups!
- We are committed to maintaining the BOOM open-source repository
- We're hiring...

A 2-person tapeout takes a village!



- RISC-V ISA
 - very out-of-order friendly!
- Chisel hardware construction language
 - object-oriented, functional programming
- FIRRTL
 - exposed RTL intermediate representation (IR)
- Rocket-chip
 - A full working SoC platform built around the Rocket in-order core
- Thanks to:
 - Krste Asanović, Rimas Avizienis, Jonathan Bachrach, Scott Beamer, David Biancolin, Christopher Celio, Henry Cook, Palmer Dabbelt, John Hauser, Adam Izraelevitz, Sagar Karandikar, Ben Keller, Donggyu Kim, Jack Koenig, Jim Lawson, Yunsup Lee, Richard Lin, Eric Love, Martin Maas, Chick Markley, Albert Magyar, Howard Mao, Miquel Moreto, Quan Nguyen, Albert Ou, David A. Patterson, Brian Richards, Colin Schmidt, Wenyu Tang, Stephen Twigg, Huy Vo, Andrew Waterman, Angie Wang, and more...



Questions?

Funding Acknowledgements

- *Research partially funded by DARPA Award Number HR0011-12-2-0016, the Center for Future Architecture Research, a member of STARnet, a Semiconductor Research Corporation program sponsored by MARCO and DARPA, and ASPIRE Lab industrial sponsors and affiliates Intel, Google, Huawei, Nokia, NVIDIA, Oracle, and Samsung.*
- *Approved for public release; distribution is unlimited. The content of this presentation does not necessarily reflect the position or the policy of the US government and no official endorsement should be inferred.*
- *Any opinions, findings, conclusions, or recommendations in this paper are solely those of the authors and does not necessarily reflect the position or the policy of the sponsors.*