


DSP ISA Extensions for an Open-Source RISC-V Implementation



PULP

Parallel Ultra Low Power



Davide Schiavone

Davide Rossi

Michael Gautschi

Eric Flamand

Andreas Traber


Luca Benini

ETH zürich

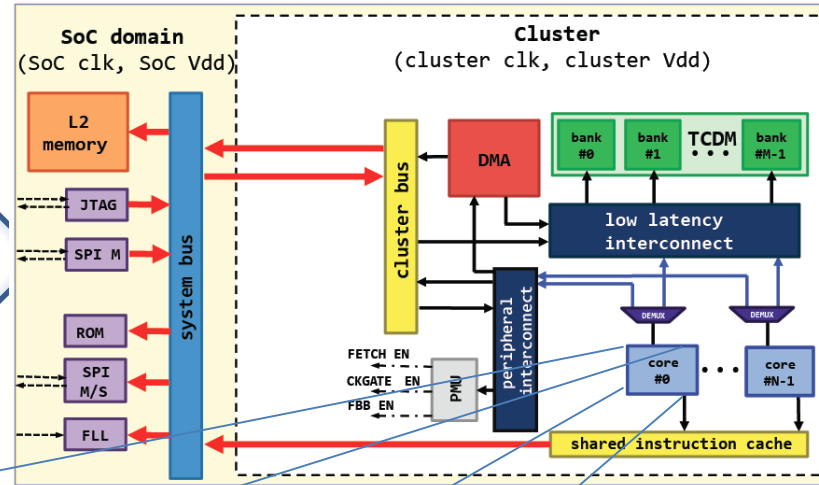
Integrated Systems Laboratory

Introduction: a typical embedded system

Sensor



Analyze and Classify



Transmit



RISCV Core

inst. fetch

ld/st unit

debug info

PULP



Energy Constrained Environment

RI5CYv2 Core

1. *RI5CYv2 is an evolution of our RISC-V implementation RI5CY, in-order 4 pipeline stages
(presented by Andreas Traber during last RISC-V Workshop)*
2. *It is a RV32IMC implementation plus some PULP specific ISA extensions to target energy efficiency*
3. *We support interrupts and exceptions and we provide a debug interface*
4. *We have GCC supports for all our extensions*
5. *Profiling and core execution trace*
6. *Our core targets tightly-coupled multi-core systems*

→ Fully featured microcontroller competitive with SOA microcontrollers

Coremarks/MHz

Core	Coremark/MHz
RV32IM	2.36
RI5CYv1	2.75
RI5CYv2	2.86
ARM Cortex M4 [1]	2.85


[1] http://infocenter.arm.com/help/topic/com.arm.doc.dai0350a/DAI0350A_coremark_benchmarking.pdf

Why DSP extensions?

1. We target signal processing algorithms
2. We need to speed up our applications to put the system in low power mode as soon as possible

How to be energy efficient?

Performance increase higher than Power increase



RI5CYv1 Core previous ISA Extensions from OR10N

```
for (i = 0; i < 101; i ++)  
  d[i] = a[i] + b[i];
```

Baseline


```
mv  x5, 0  
mv  x4, 101  
Lstart:  
lw   x2, 0(x10)  
lw   x3, 0(x11)  
addi x10, x10, 4  
addi x11, x11, 4  
add  x2, x3, x2  
sw   x2, 0(x12)  
addi x4, x4, -1  
addi x12, x12, 4  
bne  x4, x5, Lstart
```

Auto-incr load/store

```
mv  x5, 0  
mv  x4, 101  
Lstart:  
lw   x2, 0(x10!)  
lw   x3, 0(x11!)  
add  x2, x3, x2  
addi x4, x4, -1  
sw   x2, 0(x12!)  
bne  x4, x5, Lstart
```

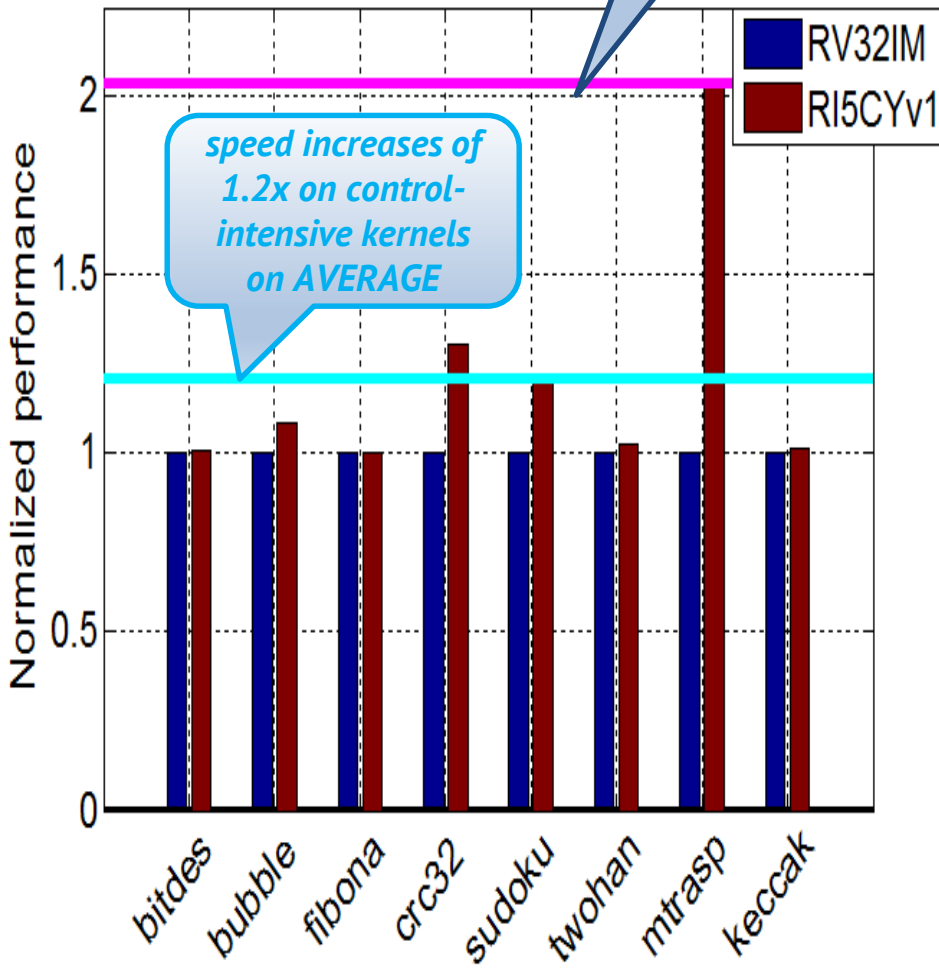
HW Loop

```
lp.setupi 101, Lend  
lw   x2, 0(x10!)  
lw   x3, 0(x11!)  
add  x2, x3, x2  
Lend: sw x2, 0(x12!)
```



RI5CYv1 vs RV32IM

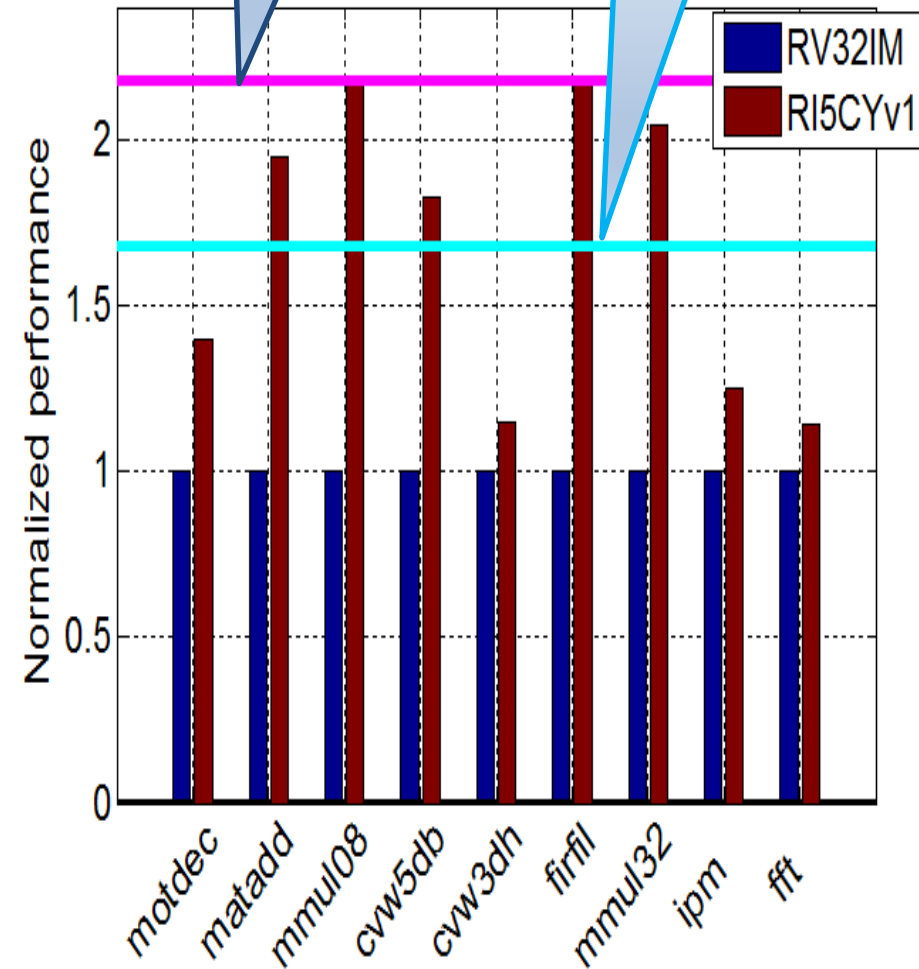
Control-Intensive kernels RI5CYv1 vs RV32IM



speed increases UP TO 2.2x on data-intensive kernels

speed increases of 1.7x on data-intensive kernels on AVERAGE

Data-Intensive kernels RI5CYv1 vs RV32IM




RI5CYv2 Core ISA Extensions

RI5CYv2 provides a set of extensions for small vectors manipulation typical of several DSP applications:

- 1. Dot product between vectors*
- 2. Saturation between $[-2^n \text{ or } 0 ; 2^n - 1]$*
- 3. Mul/Add/Sub plus round and normalization
(eg: fixed point)*
- 4. Shuffle operations for vectors*
- 5. Packed-SIMD ALU operations
(ported from our previous OpenRisc implementations)*
- 6. Bit manipulations*

Packed-SIMD
Architecture




GCC support

*R15CYv2 new instructions are supported by our GCC.
We use attributes and built-in functions in our C code*

```
typedef char charVec  
__attribute__((vector_size (4)));
```

```
charVec a = {1, 2, 3, 4};
```

```
charVec res = a + a;  
// res = {2,4,6,8}
```




```
pv.add.b x5, x4, x3  
sw x5, 0(x10)
```

```
typedef short shortVec  
__attribute__((vector_size (2)));
```

```
shortVec a = {5, 6};  
shortVec b = {-1, -2};
```

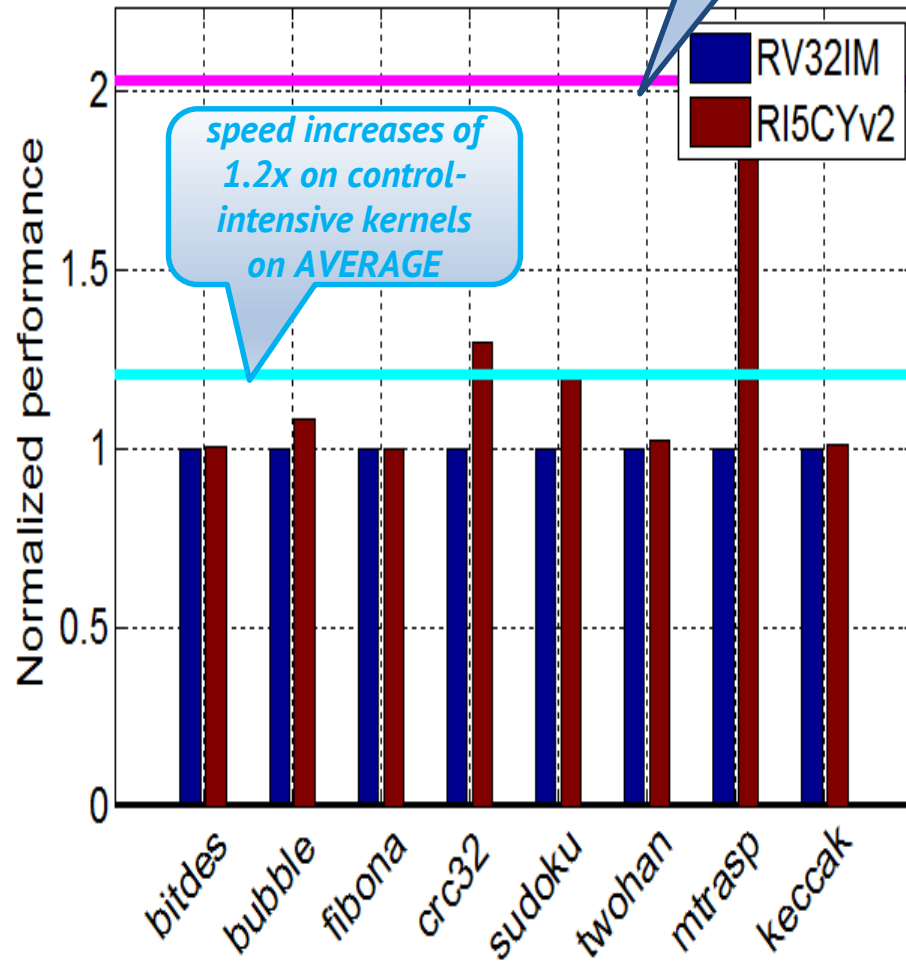
```
int res = __builtin_pulp_dotsp2(a,b);  
// res = -1*5 -2*6 = -17
```



```
pv.dotsp.h x5, x4, x3  
sw x5, 0(x10)
```

RI5CYv2 vs RV32IM

Control-Intensive kernels RI5CYv2 vs RV32IM




speed increases UP TO 2x on control-intensive kernels

speed increases UP TO 2.2x on data-intensive kernels

speed increases of 1.7x on data-intensive kernels on AVERAGE

Data-Intensive kernels RI5CYv2 vs RV32IM



Optimized C code


eg: 2D-Convolution

$$\sum_{k_1=-\infty}^{\infty} \sum_{k_2=-\infty}^{\infty} h(k_1, k_2) x(n_1 - k_1, n_2 - k_2)$$

Kernel

Load 5 vectors from the Kernel

K1, K2, K3, K4, K5



Load 5 vectors from the Image

V1, V2, V3, V4, V5

Compute the dot products

$$S = V1 * K1$$


$$S += V2 * K2$$

$$S += V3 * K3$$

$$S += V4 * K4$$

$$S += V5 * K5$$

Doing it for each output pixel



2D-Convolution with our extensions: Pseudo C Code

Scalar

```

For each column c
  For each row r
    S = 0;
    for (i = -1; i <= 1; i++) {
      for (j = -1; j <= 1; j++) {
        compute index k, w
        data = InImg[k];
        coeff = Kernel[w];
        S = S + coeff*data;
      }
    }
    S = S + ROUNDBIT;
    S = S >> (DATAWIDTH-1);
    S = S > SAT ? SAT : S;
    S = S < 0 ? 0 : S;
    OutImg[t] = S;
  
```

Vector

load the 5 vectors from Filter

For each column c

load 5 vectors from InImg[c]

For each row R

S = __builtin_pulp_dotsp2(V1,K1);

S = __builtin_pulp_sdotsp2(V2,K2, S);

S = __builtin_pulp_sdotsp2(V3,K3, S);

S = __builtin_pulp_sdotsp2(V4,K4, S);

S = __builtin_pulp_sdotsp2(V5,K5, S);

S = __builtin_pulp_addRN(S,0,DATAWIDTH-1,ROUNDBIT);

//clip inferred automatically

S = S > SAT ? SAT : S;

S = S < 0 ? 0 : S; CLIP

OutImg[t] = S; STORE PIXEL

load the next 2 new vectors V6, V7 from InImg

V1 = V2;

V2 = V3;

V3 = V6;

V4 = __builtin_pulp_shuffle2h(V4,V5,mask);

V5 = V7;

5 DOT PRODUCTS


ROUND AND SHIFT

SHIFT BACK THE WINDOW


RI5CYv2 optC vs RI5CYv1 vs RV32IM on data-intensive kernels

Using the optimized C code performance increases up to 9.5x!!!

*speed increases UP
TO 9.5 x on data-
intensive kernels*



RI5CYv2 optC vs RI5CYv1 on data-intensive kernels



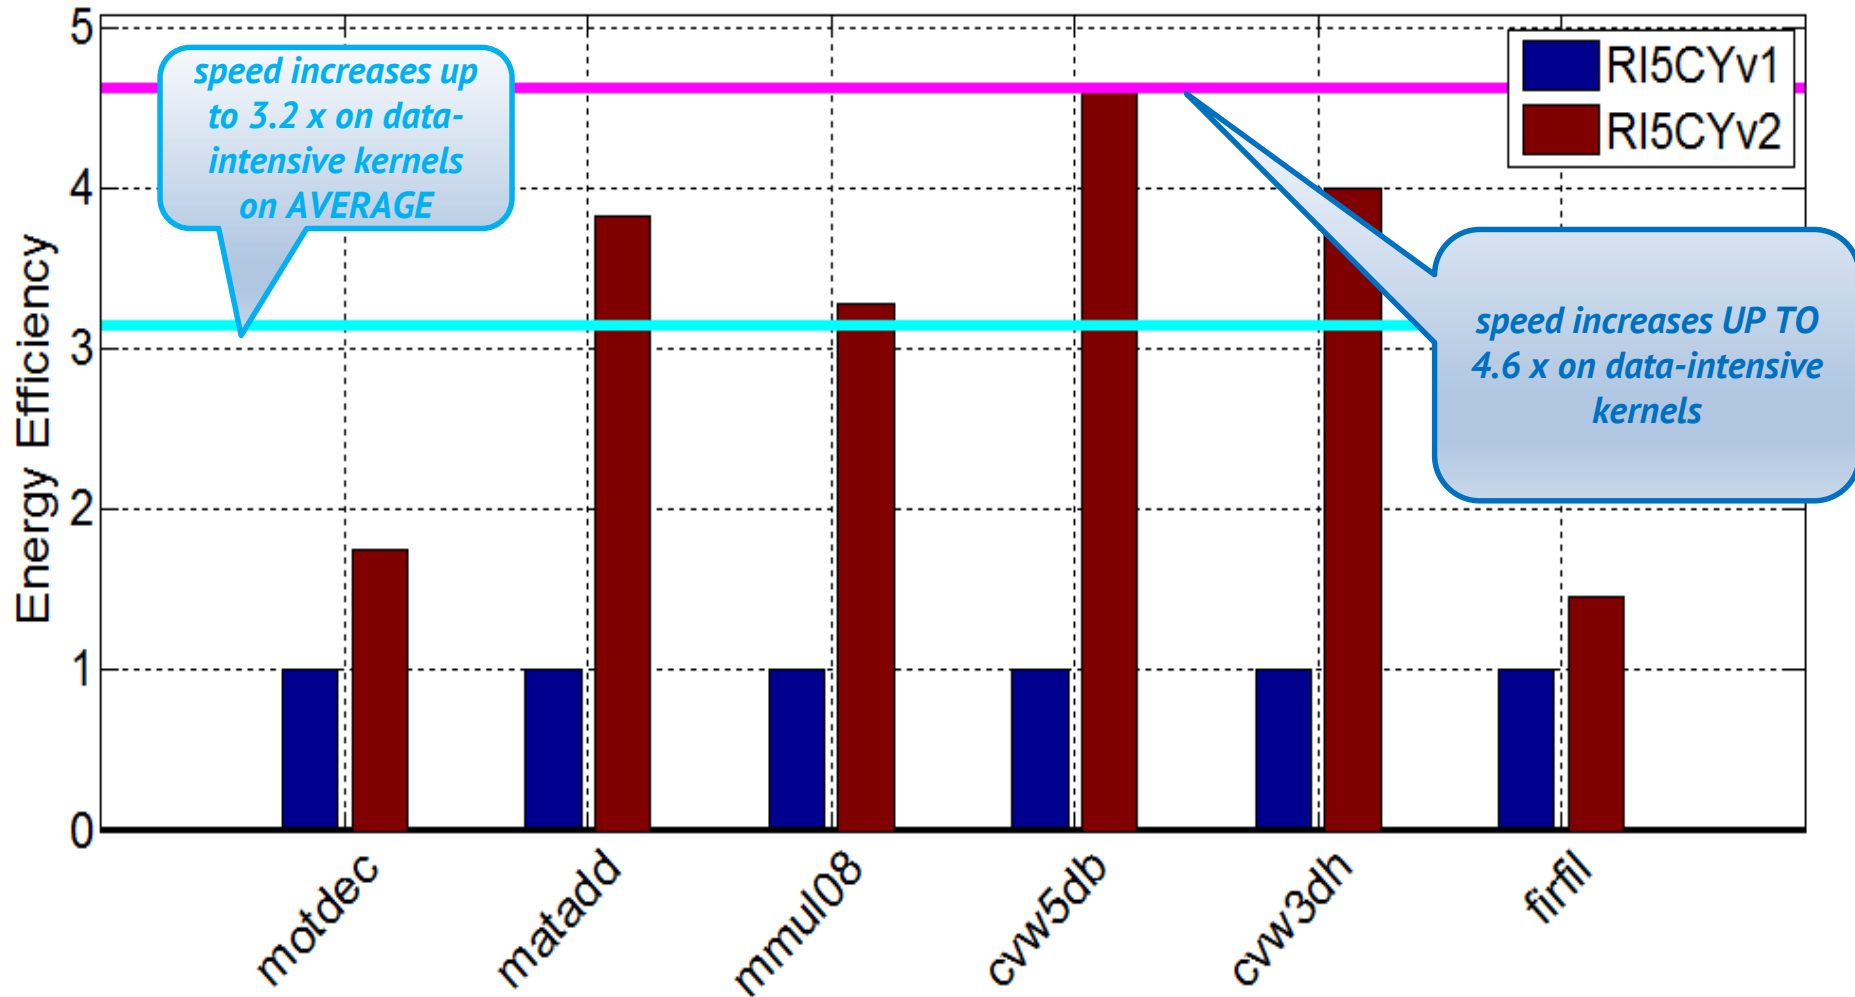
RI5CYv2: Timing, Area and Power

	RI5CYv2
Technology	65 nm
Conditions	25°C, 1.2V
Dynamic Power [$\mu W / MHz$]	30.6
Area [mm^2]	0.074
Max Frequency	650 MHz

*Running dsp optimized kernels increases power up to +14%, **but...***

Energy: RI5CYv2 optC vs RI5CYv1 on data-intensive kernels [*]


Energy Efficiency Data-Intensive kernels RI5CYv2 RI5CYv1




[*] Power estimated using the same core with different compiler options

Building PULP: silicon roadmap


- **Main PULP chips** (ST 28nm FDSOI)
 - PULPv1
 - PULPv2
 - PULPv3 (under test)
 - **PULPv4 (RISC-V based systems) (in progress)**
- **PULP development** (UMC 65nm)
 - Artemis - IEEE 754 FPU
 - Hecate - Shared FPU
 - Selene - Logarithmic Number System FPU
 - Diana - Approximate FPU
 - Mia Wallace – full system
 - **Imperio - PULPino chip (RISC-V based systems)**
 - Fulmine – Secure HW accelerator cluster
- **PULP development** (GF 28nm)
 - **Honey Bunny (RISC-V based systems)**
- **Early building blocks** (UMC180)
 - Sir10us
 - Or10n
- **Mixed-signal systems** (SMIC 130nm)
 - VivoSoC
 - VivoSoc2
- **IcySoC chips approx. computing platforms** (ALP 180nm)
 - Diego
 - Manny
 - Sid



28nm




28nm




28nm




65nm




65nm




65nm




65nm




65nm




130nm




180nm




180nm



180nm




180nm



180nm

Users and collaborations


7 COMPANIES



7 UNIVERSITIES



PULP
Parallel Ultra Low Power



*...they are growing
...and growing*

RELEASED

TO BE RELEASED



Target Platforms and SDK

1. *Our PULP systems target mainly ASIC and FPGA designs*

2. *Virtual Platform*

- *C++ for fast execution and Python for configurations*
- *ISS to emulate our RISC-V*
- *Timing model: event based*
- *DMA, Memories, Peripherals models*
- *Time accuracy ~10%*


3. *OpenMP runtime support*



released soon

Conclusions

1. Thanks to all our extensions, **typical signal processing applications** like convolutions, gain up to 9.5x with optimized C code in speed. All of our extensions are supported by our released compiler.
2. On average the extended core has a small **increase in power** and a **big increase in performance** which leads to a **higher energy efficiency**
→ **Energy efficiency increased** up to 4.6x with optimized C code with respect the general purpose implementation
1. You can download the HDL of our **RISC-V** core here
<http://www.pulp-platform.org/>
2. We are going to release **PULP** in December, meanwhile you can use our core on PULPino!
See our release plans:
<http://www.pulp-platform.org/release-plan/>




We are open to collaborate with you

We would like to receive your feedbacks

- *Help us to improve our work*
- *Please report bugs and fixes → We are looking for verification support*

Thank you!!!
QUESTIONS?



<http://www.pulp-platform.org>