

Improving the Performance Per Area Factor of RISC-V Based Multi-Core Systems

Tobias Strauch
R&D, EDaptix, Munich, Germany
tobias@edaptix.com


4th RISC-V Workshop at MIT, Boston, USA

July 13, 2016



Agenda

- ▶ C-Slow Retiming
- ▶ System Hyper Pipelining
 - ▶ Basic Idea
 - ▶ Performance Balancing
 - ▶ Deep Pipelining
 - ▶ Extended Pipelining
 - ▶ Performance per Area Factor
- ▶ microRISC Project
- ▶ miniRISC Project
- ▶ Miscellaneous



C-Slow Retiming





C-Slow Retiming





C-Slow Retiming





C-Slow Retiming





C-Slow Retiming






C-Slow Retiming







C-Slow Retiming







C-Slow Retiming







C-Slow Retiming







C-Slow Retiming







C-Slow Retiming







C-Slow Retiming




C-Slow Retiming



C-Slow Retiming



C-Slow Retiming




- ▶ Known since the 60's, Barrel processors
- ▶ Leiserson [1]
- ▶ Berkeley summer class in 2001
- ▶ CSR on FPGAs (LUT level) by Weaver et al. [2]
- ▶ Millions of engineers


C-Slow Retiming (Personal Work)

- ▶ Diploma Thesis in 1998 on CSR of FSM
- ▶ LSI Logic (Milpitas, CA), RISC Processor, CSR on gate level
- ▶ 2001: Timing estimation on RTL, RTL code modification, ...
- ▶ 2010: Automatically apply CSR on RTL
- ▶ 2010: AVR Core (VHDL) on opencores.org
- ▶ 2010: OpenRISC Core (Verilog) on opencores.org
- ▶ Papers on CSR on RTL, CSR in safety critical designs, ...


System Hyper Pipelining




$C=3, D=16$




$P=3, C=3, D=16$




$C=3$




$P=3, C=3$




$P=3$





System Hyper Pipelining




$C=3, D=16$




$C=3$



$P=3, C=3, D=16$




$P=3, C=3$




$P=3$



System Hyper Pipelining




C=3, D=16




C=3



P=3, C=3, D=16




P=3, C=3




P=3


System Hyper Pipelining




C=3, D=16




P=3, C=3, D=16




C=3




P=3, C=3




P=3





System Hyper Pipelining




$C=3, D=16$




$C=3$



$P=3, C=3, D=16$




$P=3, C=3$




$P=3$



System Hyper Pipelining




$C=3, D=16$




$C=3$



$P=3, C=3, D=16$




$P=3, C=3$




$P=3$


System Hyper Pipelining




$C=3, D=16$




$P=3, C=3, D=16$




$C=3$




$P=3, C=3$




$P=3$





System Hyper Pipelining




$C=3, D=16$




$C=3$



$P=3, C=3, D=16$




$P=3, C=3$




$P=3$


System Hyper Pipelining




$C=3, D=16$




$P=3, C=3, D=16$




$C=3$




$P=3, C=3$




$P=3$




System Hyper Pipelining




C=3, D=16




P=3, C=3, D=16




C=3




P=3, C=3




P=3





System Hyper Pipelining




$C=3, D=16$




$C=3$



$P=3, C=3, D=16$




$P=3, C=3$




$P=3$


System Hyper Pipelining




$C=3, D=16$




$P=3, C=3, D=16$




$C=3$




$P=3, C=3$




$P=3$




System Hyper Pipelining




$C=3, D=16$




$P=3, C=3, D=16$




$C=3$



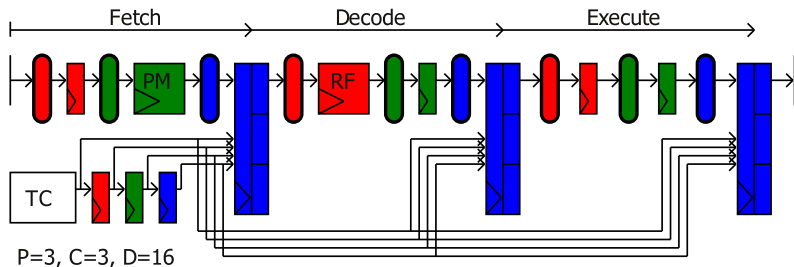
$P=3, C=3$



$P=3$




System Hyper Pipelining




- ▶ Based on CSR
- ▶ Replaces original registers with memories (D).
- ▶ Adds thread stalling and therefore thread bypassing features.
- ▶ Cycle accurate performance balancing
- ▶ Late read / early write [3]

SHP, Performance Balancing




- ▶ Predictive runtime of specific threads
- ▶ Sync. start/execution, "Dynamic Length Instruction Words"

SHP, Deep Pipelining, T0 in “Beast Mode”



- ▶ Detect instruction dependency, LSB RISC-V ISA
- ▶ 1) Instruction LSB-sniffing
- ▶ 2) Enhanced PM read
- ▶ 3) Consider stall/flush signal
- ▶ No back-to-back, turned on/off on the fly

SHP, Extended Pipelining, “Frequency-Over-Scaling”




- ▶ Less registers per path than necessary
- ▶ Multi-cycle path for example in datapath section
- ▶ Re-execution of thread with valid multi-cycle path



SHP, System Level

- ▶ Performance (frequency) over area curve
- ▶ System level performance improvements !!!
- ▶ microRISC (Vscale), microcontroller, virtual peripherals, ...
- ▶ miniRISC (lowRISC, Rocket), SoC, OS, FPU, accelerators, ...

microRISC



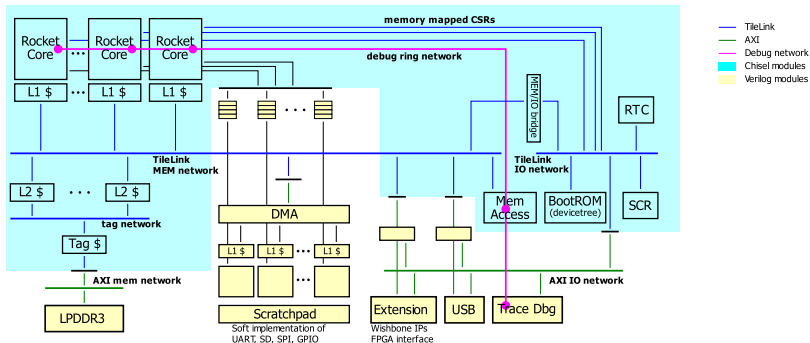
- ▶ Vscale @80MHz async DDR3, SHP-ed @250MHz synchron
- ▶ Based on RV32IM subset (no FENCE and no SI, "RV32B")
- ▶ Individual threads handle interrupts (less stack activity).
- ▶ Main focus on virtual peripherals
- ▶ TC with event handler (complex timer)



- ▶ UARTs
- ▶ PWMs
- ▶ ...



- ▶ Standard toolchain compatible (C++, ...)
- ▶ Demo: How to run timing critical software peripherals based on a highly dynamic SHP-ed core.

lowRISC / University of Cambridge



- ▶ Diagram: Wei Song, lowRISC / University of Cambridge
- ▶ Wei: "personal view"

miniRISC (based on lowRISC)



- ▶ Minions part of Rocket core, heterogeneous multi-core
- ▶ SHP impact on performance per area and SoC architecture
- ▶ Arbiter/multiplexer: blocking
- ▶ Multilayer: complex
- ▶ SHP: time sliced usage at higher speed

Miscellaneous

- ▶ Simple programming model
- ▶ Fork-join operations, OpenMP
- ▶ Estimated ASIC numbers in paper [3]
- ▶ Altera Hyper Pipelining technology looks promising
- ▶ More on-FPGA memory (memory wall)
- ▶ Power consumption (work ongoing)
- ▶ Source code of projects in PDVL (VHDL, Verilog)

References

- [1] C. Leiserson and J. Saxe, “Retiming Synchronous Circuitry“, Algorithmica, vol. 6, no. 1, pp. 5-35, 1991.

- [2] N. Weaver, Y. Markovskiy, Y. Patel and J. Wawrzynek, “Post-Placement C-slow Retiming for the Xilinx Virtex FPGA,“ FPGA 2003, February 23-25, 2003, Monterey, CA, USA.

- [3] T. Strauch, “The Effects of System Hyper Pipelining on Three Computational Benchmarks Using FPGAs“, 11th International Symposium in Applied Reconfigurable Computing, ARC 2015, 13-17 April 2015, Bochum, Germany, pp. 280-290.

You made it !!!

- ▶ Thank you
- ▶ @arduissimo
- ▶ www.cloudx.cc
- ▶ tobias@cloudx.cc
- ▶ Call for cooperation: SHP-ed CPU in an ASIC technology